

Heart Disease Prediction

Peter Thai

25 April 2025

Introduction

This project applies supervised machine learning techniques to predict the presence of heart disease using the UCI Heart Disease dataset. We will explore the data, train models, and evaluate performance based on accuracy and AUC metrics.

Load Libraries and Data

```
if (!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")

## Loading required package: tidyverse

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.4.4      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

if (!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")

## Loading required package: caret
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
## lift

if (!require(randomForest)) install.packages("randomForest", repos = "http://cran.us.r-project.org")
```

```
## Loading required package: randomForest
## randomForest 4.7-1.1
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
##
## The following object is masked from 'package:dplyr':
##
##   combine
##
## The following object is masked from 'package:ggplot2':
##
##   margin

if (!require(pROC)) install.packages("pROC", repos = "http://cran.us.r-project.org")

## Loading required package: pROC
## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##   cov, smooth, var

library(tidyverse)
library(caret)
library(randomForest)
library(pROC)

heart <- read.csv("heart.csv")
heart$target <- as.factor(heart$target)
```

Exploratory Data Analysis

```
summary(heart)
```

```
##      age      sex      cp      trestbps
##  Min.   :29.00  Min.   :0.0000  Min.   :0.000  Min.    : 94.0
##  1st Qu.:47.50  1st Qu.:0.0000  1st Qu.:0.000  1st Qu.:120.0
##  Median :55.00  Median :1.0000  Median :1.000  Median :130.0
##  Mean   :54.37  Mean   :0.6832  Mean   :0.967  Mean   :131.6
##  3rd Qu.:61.00  3rd Qu.:1.0000  3rd Qu.:2.000  3rd Qu.:140.0
##  Max.   :77.00  Max.   :1.0000  Max.   :3.000  Max.   :200.0
##      chol      fbs      restecg      thalach
##  Min.   :126.0  Min.   :0.0000  Min.   :0.0000  Min.    : 71.0
##  1st Qu.:211.0  1st Qu.:0.0000  1st Qu.:0.0000  1st Qu.:133.5
##  Median :240.0  Median :0.0000  Median :1.0000  Median :153.0
##  Mean   :246.3  Mean   :0.1485  Mean   :0.5281  Mean   :149.6
##  3rd Qu.:274.5  3rd Qu.:0.0000  3rd Qu.:1.0000  3rd Qu.:166.0
```

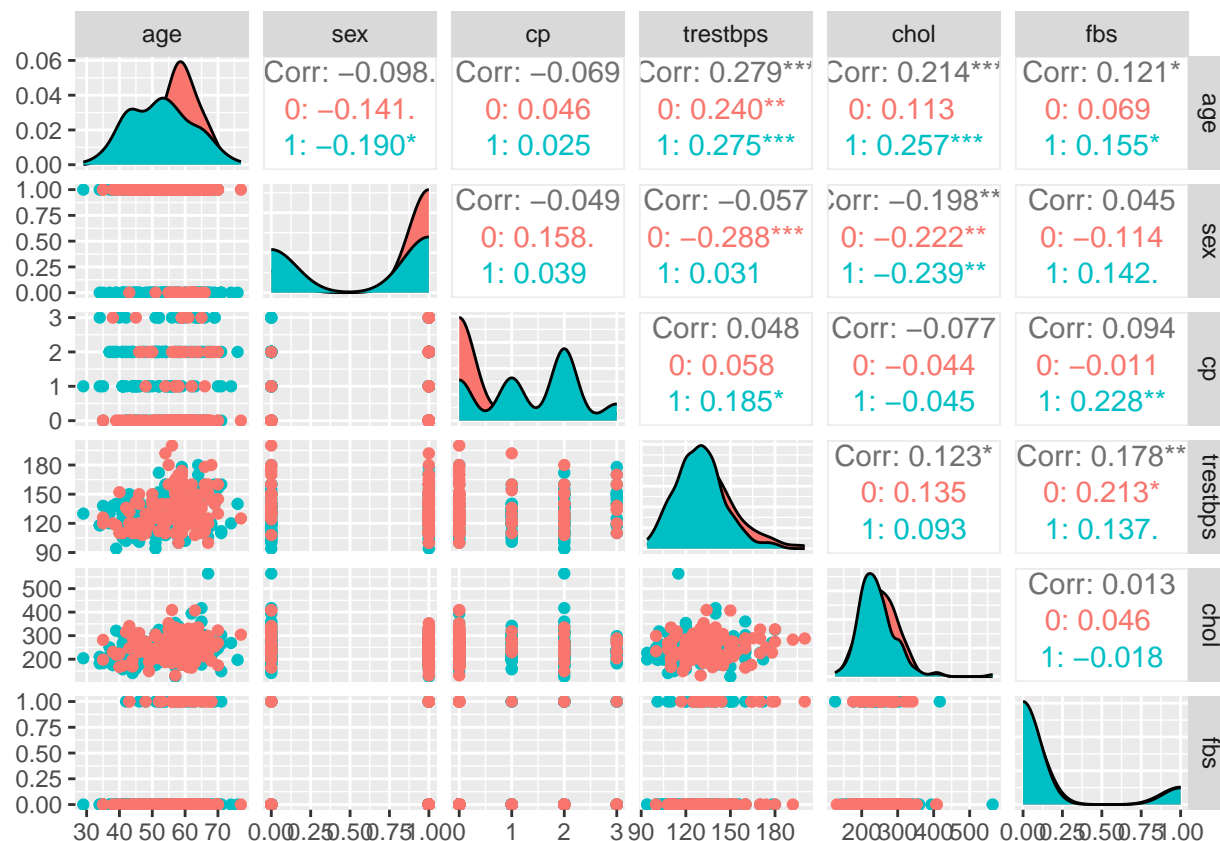
```
## Max. :564.0 Max. :1.0000 Max. :2.0000 Max. :202.0
## exang oldpeak slope ca
## Min. :0.0000 Min. :0.00 Min. :0.000 Min. :0.0000
## 1st Qu.:0.0000 1st Qu.:0.00 1st Qu.:1.000 1st Qu.:0.0000
## Median :0.0000 Median :0.80 Median :1.000 Median :0.0000
## Mean :0.3267 Mean :1.04 Mean :1.399 Mean :0.7294
## 3rd Qu.:1.0000 3rd Qu.:1.60 3rd Qu.:2.000 3rd Qu.:1.0000
## Max. :1.0000 Max. :6.20 Max. :2.000 Max. :4.0000
## thal target
## Min. :0.000 0:138
## 1st Qu.:2.000 1:165
## Median :2.000
## Mean :2.314
## 3rd Qu.:3.000
## Max. :3.000
```

```
str(heart)
```

```
## 'data.frame': 303 obs. of 14 variables:
## $ age : int 63 37 41 56 57 57 56 44 52 57 ...
## $ sex : int 1 1 0 1 0 1 0 1 1 1 ...
## $ cp : int 3 2 1 1 0 0 1 1 2 2 ...
## $ trestbps: int 145 130 130 120 120 140 140 120 172 150 ...
## $ chol : int 233 250 204 236 354 192 294 263 199 168 ...
## $ fbs : int 1 0 0 0 0 0 0 0 1 0 ...
## $ restecg : int 0 1 0 1 1 1 0 1 1 1 ...
## $ thalach : int 150 187 172 178 163 148 153 173 162 174 ...
## $ exang : int 0 0 0 0 1 0 0 0 0 0 ...
## $ oldpeak : num 2.3 3.5 1.4 0.8 0.6 0.4 1.3 0 0.5 1.6 ...
## $ slope : int 0 0 2 2 2 1 1 2 2 2 ...
## $ ca : int 0 0 0 0 0 0 0 0 0 0 ...
## $ thal : int 1 2 2 2 2 1 2 3 3 2 ...
## $ target : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 2 2 2 ...
```

```
heart %>% GGally::ggpairs(columns = 1:6, aes(color = target))
```

```
## Registered S3 method overwritten by 'GGally':
## method from
## +.gg ggplot2
```



Data Partitioning

```
set.seed(123)
train_index <- createDataPartition(heart$target, p = 0.8, list = FALSE)
train_data <- heart[train_index, ]
test_data <- heart[-train_index, ]
```

Logistic Regression Model

```
log_model <- glm(target ~ ., data = train_data, family = "binomial")
summary(log_model)
```

```
##
## Call:
## glm(formula = target ~ ., family = "binomial", data = train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.654796   2.801571   0.234  0.815199
## age         0.016291   0.026641   0.612  0.540861
```

```
## sex          -1.973555   0.546594  -3.611 0.000305 ***
## cp           0.796468   0.204425   3.896 9.77e-05 ***
## trestbps     -0.013829   0.011577  -1.195 0.232282
## chol         -0.004440   0.004294  -1.034 0.301066
## fbs          0.282036   0.612297   0.461 0.645071
## restecg      0.895128   0.408819   2.190 0.028557 *
## thalach      0.035307   0.011729   3.010 0.002609 **
## exang        -0.968394   0.466916  -2.074 0.038077 *
## oldpeak      -0.638153   0.247076  -2.583 0.009800 **
## slope        0.305689   0.415728   0.735 0.462151
## ca           -0.765629   0.211001  -3.629 0.000285 ***
## thal         -1.129838   0.340057  -3.322 0.000892 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 335.05  on 242  degrees of freedom
## Residual deviance: 161.07  on 229  degrees of freedom
## AIC: 189.07
##
## Number of Fisher Scoring iterations: 6
```

```
log_probs <- predict(log_model, test_data, type = "response")
log_preds <- ifelse(log_probs > 0.5, 1, 0) %>% as.factor()
confusionMatrix(log_preds, test_data$target)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0  1
##           0 19  3
##           1  8 30
##
##              Accuracy : 0.8167
##              95% CI : (0.6956, 0.9048)
##    No Information Rate : 0.55
##    P-Value [Acc > NIR] : 1.344e-05
##
##              Kappa : 0.6233
##
## Mcnemar's Test P-Value : 0.2278
##
##              Sensitivity : 0.7037
##              Specificity : 0.9091
##              Pos Pred Value : 0.8636
##              Neg Pred Value : 0.7895
##              Prevalence : 0.4500
##              Detection Rate : 0.3167
##              Detection Prevalence : 0.3667
##              Balanced Accuracy : 0.8064
##
##              'Positive' Class : 0
##
```


ROC Curve - Logistic Regression

```
log_roc <- roc(test_data$target, log_probs)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(log_roc, main = "ROC Curve - Logistic Regression")
```



```
auc(log_roc)
```

```
## Area under the curve: 0.8676
```

Random Forest Model

```
set.seed(123)
rf_model <- randomForest(target ~ ., data = train_data, ntree = 500, mtry = 4, importance = TRUE)
print(rf_model)
```

```
##
## Call:
## randomForest(formula = target ~ ., data = train_data, ntree = 500,      mtry = 4, importance = TRUE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 4
##
##           OOB estimate of  error rate: 19.34%
## Confusion matrix:
##      0   1 class.error
## 0 84  27   0.2432432
## 1 20 112   0.1515152
```

```
rf_preds <- predict(rf_model, test_data)
confusionMatrix(rf_preds, test_data$target)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 20   5
##           1   7 28
##
##           Accuracy : 0.8
##           95% CI : (0.6767, 0.8922)
##           No Information Rate : 0.55
##           P-Value [Acc > NIR] : 4.67e-05
##
##           Kappa : 0.5932
##
##           McNemar's Test P-Value : 0.7728
##
##           Sensitivity : 0.7407
##           Specificity : 0.8485
##           Pos Pred Value : 0.8000
##           Neg Pred Value : 0.8000
##           Prevalence : 0.4500
##           Detection Rate : 0.3333
##           Detection Prevalence : 0.4167
##           Balanced Accuracy : 0.7946
##
##           'Positive' Class : 0
##
```


ROC Curve - Random Forest

```
rf_probs <- predict(rf_model, test_data, type = "prob")[,2]
rf_roc <- roc(test_data$target, rf_probs)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(rf_roc, main = "ROC Curve - Random Forest")
```



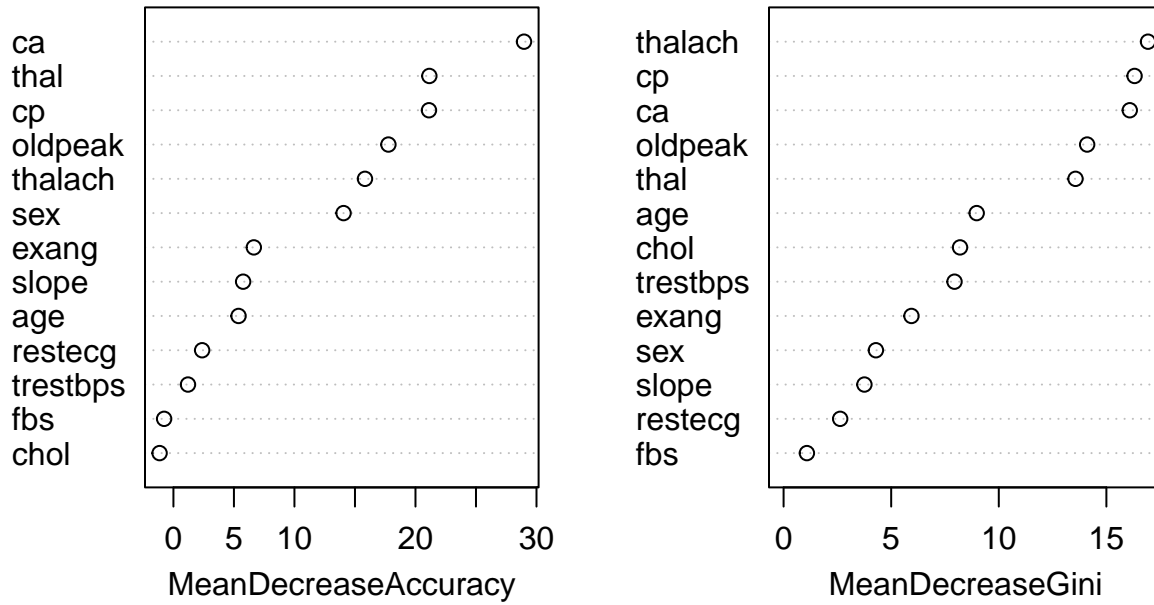
```
auc(rf_roc)
```

```
## Area under the curve: 0.8782
```

Feature Importance

```
varImpPlot(rf_model)
```


rf_model



Conclusion

Both models performed well, with Random Forest showing slightly higher predictive power based on AUC. This demonstrates the value of ensemble methods for structured healthcare data.

References

- UCI Machine Learning Repository: Heart Disease Dataset
- caret, randomForest, and pROC R packages